


Recitation Exercises:

1.1 Chapter 8

1.1.1 Exercise 1



1.1.2 Exercise 3

```

p1 = seq(0 + 1e-06, 1 - 1e-06, length.out = 100)
p2 = 1 - p1


class_error = 1 - apply(rbind(p1, p2), 2, max)

gini_index = p1 * (1 - p1) + p2 * (1 - p2)



cross_entropy = -(p1 * log(p1) + p2 * log(p2))

plot(p1, class_error, type = "l", col = "black", xlab = "Pm1", ylab = "Error
metrics", ylim = c(min(c(class_error, gini_index, cross_entropy)), max(class_error,
gini_index, cross_entropy)))
lines(p1, gini_index, col = "blue")
lines(p1, cross_entropy, col = "red")
legend(0.3, 0.2, c("Classification error", "Gini index", "Cross entropy"), col =
c("black", "blue", "red"), lty = c(1, 1))
grid()

```



1.1.3 Exercise 4



1.1.4 Exercise 5

- Using Majority approach, X can be classified to the **Red Class**, as we have 6 out of 10 probabilities for Class Red (0.55, 0.6, 0.6, 0.65, 0.7, 0.75).
- Using Average probability approach, the average probability comes out to be, $(0.1 + 0.15 + 0.2 + 0.2 + 0.55 + 0.6 + 0.6 + 0.65 + 0.7 + 0.75) / 10 = 0.45$. Thus, X can be classified to the **Green Class**.

1.2 Chapter 10

1.2.1 Exercise 1

a. Proof:

$$\text{Given: } \bar{x}_{kj} = \frac{1}{|C_k|} \sum_{i \in C_k} x_{ij}$$

$$\text{LHS} = \frac{1}{|C_k|} \sum_{i,i' \in C_k} \sum_{j=1}^p (x_{ij} - \bar{x}_{kj})^2$$

$$= \frac{1}{|C_k|} \sum_{i,i' \in C_k} \sum_{j=1}^p (x_{ij} - x_{i'j})^2 \quad \dots \text{Expanding as } (a - b)^2 = a^2 - 2ab + b^2$$

$$= \sum_i \sum_j x_{ij}^2 - 2 \sum_i \sum_j x_{ij} \bar{x}_{kj} + \sum_i \sum_j x_{ij}^2$$

$$= 2 \sum_i \sum_j x_{ij}^2 - 2 |C_k| \sum_j \bar{x}_{kj}^2 \quad \dots (1)$$

$$\text{RHS} = 2 \sum_{i \in C_k} \sum_{j=1}^p (x_{ij} - \bar{x}_{kj})^2$$

$$= 2 \sum_i \sum_j x_{ij}^2 - 4 \sum_i \sum_j x_{ij} \bar{x}_{kj} + 2 \sum_i \sum_j \bar{x}_{kj}^2 \quad \dots \text{Expanding as } (a - b)^2 = a^2 - 2ab + b^2$$

$$= 2 \sum_i \sum_j x_{ij}^2 - 2 |C_k| \sum_j \bar{x}_{kj}^2 \quad \dots (2)$$

As from equations (1) and (2), LHS = RHS

Hence Proved.

- b. Objective 10.11 can be seen at the step 2.b of Algorithm 10.1. It says, when we assign a new observation to the cluster whose centroid is closest, we decrease the right member of the identity. So, we decrease the left member of the identity which is our objective. Overall the identity 10.11 shows that minimizing the sum of the squared Euclidean distance for each cluster is the same as minimizing the within-cluster variance for each cluster.

1.2.2 Exercise 2

Following calculations describes how Single and Complete linkage is performed:

	1	2	3	4
1	0	0.3	0.4	0.7
2	0.3	0	0.5	0.8
3	0.4	0.5	0	0.45
4	0.7	0.8	0.45	0

Complete Linkage

smallest inter-cluster
dist = 0.3, clustering
point 1 and 2

[1,2]	3	4	
1,2	0	0.5	0.8
3	0.5	0	0.45
4	0.8	0.45	0

$$\rightarrow d([1,2], 3)$$

$$= \max[d(1,3), d(2,3)]$$

$$= \max[0.4, 0.5] = \underline{\underline{0.5}}$$

$$\rightarrow d([1,2], 4)$$

$$= \max[d(1,4), d(2,4)]$$

$$= \max[0.7, 0.8] = \underline{\underline{0.8}}$$

Single Linkage

smallest inter-cluster
dis = 0.3, clustering pair
1 and 2

[1,2]	3	4	
1,2	0	0.4	0.7
3	0.4	0	0.45
4	0.7	0.45	0

$$\rightarrow d([1,2], 3)$$

$$= \min[d(1,3), d(2,3)]$$

$$= \min[0.4, 0.5] = \underline{\underline{0.4}}$$

$$\rightarrow d([1,2], 4)$$

$$= \min[d(1,4), d(2,4)]$$


$$= \min[0.7, 0.8] = \underline{\underline{0.7}}$$

a.

```
dm = as.dist(matrix(c(0, 0.3, 0.4, 0.7,
                     0.3, 0, 0.5, 0.8,
                     0.4, 0.5, 0.0, 0.45,
                     0.7, 0.8, 0.45, 0.0), nrow = 4))
```

```
# Complete Linkage
plot(hclust(dm, method = "complete"))
```

Cluster Dendrogram




dm
hclust(*, "complete")

b.

```
# Single Linkage  
plot(hclust(dm, method = "single"))
```

Cluster Dendrogram




dm
hclust (*, "single")

- c. In Cluster 1: (1,2) and in Cluster 2: (3,4)
- d. In Cluster 1: (4) and in Cluster 2: (3,(1,2))
- e.

```
# Swapping the positions  
plot(hclust(dm, method = "complete"), labels = c(2,1,4,3))
```

Cluster Dendrogram




dm
hclust (*, "complete")

1.2.3 Exercise 3

a.

```
# Plotting the Observations  
x = cbind(c(1, 1, 0, 5, 6, 4), c(4, 3, 4, 1, 2, 0))  
plot(x[,1], x[,2], xlab = "X1", ylab = "X2")
```



b.

```
# Randomly assigning a cluster to each observation
set.seed(1)
labels = sample(2, nrow(x), replace = T)
labels

## [1] 1 2 1 1 2 1

plot(x[, 1], x[, 2], col = (labels + 1), pch = 20, cex = 2, xlab = "X1", ylab = "X2")
```


c. Centroid for Green Cluster

$$\begin{aligned} &= \bar{x}_{11} = \frac{1}{3} (0 + 4 + 5) = 3 \\ &= \bar{x}_{12} = \frac{1}{3} (4 + 0 + 1) = \frac{5}{3} = 1.67 \\ &\text{Centroid}_{\text{green}} = (3, 1.67) \end{aligned}$$

Centroid for Red Cluster


$$\begin{aligned} &= \bar{x}_{21} = \frac{1}{3} (1 + 1 + 6) = \frac{8}{3} = 2.67 \\ &= \bar{x}_{22} = \frac{1}{3} (2 + 4 + 3) = 3 \\ &\text{Centroid}_{\text{red}} = (2.67, 3) \end{aligned}$$

```
# Computing Centroid for each cluster
centroid1 = c(mean(x[labels == 1, 1]), mean(x[labels == 1, 2]))
centroid2 = c(mean(x[labels == 2, 1]), mean(x[labels == 2, 2]))
plot(x[,1], x[,2], col=(labels + 1), pch = 20, cex = 2, xlab = "X1", ylab = "X2")
points(centroid1[1], centroid1[2], col = 2, pch = 4)
points(centroid2[1], centroid2[2], col = 3, pch = 4)
```




d.

```
# Assigning each observation to the centroid to which it is closest, in terms of Euclidean distance
labels = c(1, 1, 1, 2, 2, 2)
plot(x[, 1], x[, 2], col = (labels + 1), pch = 20, cex = 2, xlab = "X1", ylab = "X2")
points(centroid1[1], centroid1[2], col = 2, pch = 4)
points(centroid2[1], centroid2[2], col = 3, pch = 4)
```




e.

```
# Repeating Steps (c) and (d) until the answers obtained stop changing
centroid1 = c(mean(x[labels == 1, 1]), mean(x[labels == 1, 2]))
centroid2 = c(mean(x[labels == 2, 1]), mean(x[labels == 2, 2]))
plot(x[,1], x[,2], col=(labels + 1), pch = 20, cex = 2, xlab = "X1", ylab = "X2")
points(centroid1[1], centroid1[2], col = 2, pch = 4)
points(centroid2[1], centroid2[2], col = 3, pch = 4)
```



f.

```
# Coloring the observations according to the clusters Labels obtained
plot(x[, 1], x[, 2], col=(labels + 1), pch = 20, cex = 2, xlab = "X1", ylab = "X2")
```



1.2.4 Exercise 4

- a. Given clusters: {1,2,3} and {4,5}

Complete Linkage:

$$\begin{aligned}
 &= d[\{1,2,3\}, \{4,5\}] \\
 &= \max [d(1,4), d(1,5), d(2,4), d(2,5), d(3,4), d(3,5)] \quad \dots (1)
 \end{aligned}$$

Single Linkage:

$$\begin{aligned} &= d[\{1,2,3\}, \{4,5\}] \\ &= \min [d(1,4), d(1,5), d(2,4), d(2,5), d(3,4), d(3,5)] \end{aligned} \quad \dots (2)$$

Now from (1) and (2) we have following two cases:

Case 1: If All the distances are **not** equal

So, situation (1) will give us a bigger value as it picks the maximum value and situation (2) will give us a smaller value as it picks the minimum value.

Hence, in this case **Complete Linkage will fuse at a higher height.**

Case 2: If All the distances are equal

As all the distances are equal both situation (1) and (2) will give us the same value.

Hence, in this case both **Complete and Single Linkage will fuse at the same height.**

- b. Given clusters: {5} and {6}

As there are only single points in the clusters the dissimilarity will be the same for both Single and Complete Linkage and thus **both will fuse at the same height.**

1.2.5 Exercise 6

- The 1st Principal Component explains 10% of the variation, which means, only 10% of the total information in the gene data is preserved, while the remaining 90% is lost by projecting the tissue sample observations onto the 1st Principal Component. This can also be stated as, the remaining 90% of variance in gene data is not explained by first principal component.
- Given the flaw shown in pre-analysis of a time-wise linear trend amongst the tissue samples' first principal component, I would suggest to include the machine used (A vs B) as a feature of the data set. This should enhance the PVE of the first principal component before applying the two-sample t-test.
- c.

```
set.seed(16)
Control = matrix(rnorm(50 * 1000), ncol = 50)
Treatment = matrix(rnorm(50 * 1000), ncol = 50)
X = cbind(Control, Treatment)
X[1, ] = seq(-18, 18 - .36, .36) # Linear trend in one dimension
pr.out = prcomp(scale(X))
summary(pr.out)$importance[, 1]

##      Standard deviation Proportion of Variance Cumulative Proportion
##            3.162982           0.100040           0.100040
```

We have 10.004% variance explained by the first principal component. Now, adding in A vs B via 10 vs 0 encoding.

```

X = rbind(X, c(rep(10, 50), rep(0, 50)))
pr.out = prcomp(scale(X))
summary(pr.out)$importance[, 1]

##      Standard deviation Proportion of Variance Cumulative Proportion
##            3.409353           0.116240          0.116240

```

Now we have 11.624% variance explained by the first principal component. That's an improvement of 1.62%.

2 Practicum Problems:

2.1 Problem 1

```

library(rpart)
library(rpart.plot)

gini = function(p) {
  gini.index = 2 * p * (1 - p)
  return (gini.index)
}

entropy = function(p) {
  entropy = (p * log(p) + (1 - p) * log(1 - p))
  return (entropy)
}

df1 = data.frame(x1 = rnorm(150, 5, 2), class = rep("Y", 150))
df2 = data.frame(x1 = rnorm(150, -5, 2), class = rep("N", 150))
dataset1 = rbind(df1, df2)


fit1 = rpart(class ~ x1, method="class", data = dataset1)
printcp(fit1) # cp = cost complexity parameter in CART

##
## Classification tree:
## rpart(formula = class ~ x1, data = dataset1, method = "class")
##
## Variables actually used in tree construction:
## [1] x1
##
## Root node error: 150/300 = 0.5
##
## n= 300
##
##      CP nsplit rel error    xerror     xstd
## 1 0.99333      0 1.0000000 1.120000 0.0573178
## 2 0.01000      1 0.0066667 0.013333 0.0093966

rpart.plot(fit1, main = "Classification Tree for dataset1")

```

Classification Tree for dataset1



In this case the threshold value for 1st split is 0.23. and the tree has 3 nodes 1 Root and 2 leaf. As in this case our dataset is created using normal distribution where 150 data samples are created with mean = 5 and standard deviation = 2 and 150 data sample are created with mean = -5 and standard deviation = 2 we get exact 50-50% class distribution.

```

p = c(.50, 0.01, 1)
gini_values = sapply(p, gini)
cat("Gini Values for dataset1: ",gini_values)

## Gini Values for dataset1:  0.5 0.0198 0


entropy_values = sapply(p, entropy)
cat("Entropy Values for dataset1: ",entropy_values)

## Entropy Values for dataset1: -0.6931472 -0.05600153 NaN

df3 = data.frame(x1 = rnorm(250,1,2),class = rep("Y",250))
df4 = data.frame(x1 = rnorm(250,-1,2),class = rep("N",250))
dataset2 = rbind(df3,df4)

fit2 = rpart(class ~ x1, method="class", data = dataset2)
rpart.plot(fit2, main = "Classification Tree for dataset2")
  
```

Classification Tree for dataset2



This tree has 21 nodes and on pruning we get 3 nodes. Here our dataset is created with rnorm values of (1,2) and (-1,2) against the (5,2),(-5,2) in the previous case, we are getting poor splits as compared to the previous results. Here, a greater number of nodes are created because there are more values in a smaller range, which leads to overlapping thus, a greater number of split nodes are created. This may lead to overfitting when unseen data appears and hence we prune to tree to avoid this problem.

```

p = c(.50, 0.38,
0.49,0.53,0.51,0.54,0.51,0.55,0.51,0.56,0.22,0.34,0.27,0.27,0.38,0.44,0.68,0.80,0.68,0.71
,0.81)
gini_values = sapply(p, gini)
cat("Gini Values for dataset2: ",gini_values)


## Gini Values for dataset2:  0.5 0.4712 0.4998 0.4982 0.4998 0.4968 0.4998 0.495 0.4998
0.4928 0.3432 0.4488 0.3942 0.3942 0.4712 0.4928 0.4352 0.32 0.4352 0.4118 0.3078

entropy_values = sapply(p, entropy)
cat("Entropy Values for dataset2: ",entropy_values)

## Entropy Values for dataset2: -0.6931472 -0.6640641 -0.6929472 -0.6913461 -0.6929472 -
0.6899438 -0.6929472 -0.6881388 -0.6929472 -0.6859298 -0.526908 -0.6410355 -0.5832588 -
0.5832588 -0.6640641 -0.6859298 -0.6268695 -0.5004024 -0.6268695 -0.6021517 -0.486223

# Pruning
fit2_prune = prune.rpart(fit2, cp = 0.1)
rpart.plot(fit2_prune, main = "Classification Tree for dataset2 after Pruning")
  
```

Classification Tree for dataset2 after Pruning



On Pruning the previous tree with complexity parameter, we get this above tree which has 3 nodes.

2.2 Problem 2

```

library(readr)
library(data.table)
library(corrplot)

## corrplot 0.84 loaded

wn_URL = "https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.data"
wn_data = fread(wn_URL, header = FALSE)
wn_Header = c("class", "Alcohol", "Malic_Acid", "Ash", "Alcalinity", "Magnesium",
  "Total_Phenols", "Flavanoids", "Nonflavanoid", "Proanthocyanins", "Color_Intensity", "Hue",
  "OD280/OD315", "Proline")
colnames(wn_data) = wn_Header

apply(wn_data[, -1], 2, mean)

##          Alcohol      Malic_Acid          Ash      Alcalinity
## 13.0006180     2.3363483    2.3665169   19.4949438
##      Magnesium Total_Phenols Flavanoids Nonflavanoid
## 99.7415730     2.2951124    2.0292697    0.3618539
## Proanthocyanins Color_Intensity           Hue OD280/OD315
##          1.5908989      5.0580899    0.9574494    2.6116854
##          Proline
##        746.8932584
  
```

As the means of the attributes differ, we need to apply Scaling.

```

wn_pca = prcomp(wn_data[,-1], scale = TRUE, center = TRUE)
summary(wn_pca)

## Importance of components:
##                               PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation     2.169  1.5802  1.2025  0.95863  0.92370  0.80103  0.74231
## Proportion of Variance 0.362  0.1921  0.1112  0.07069  0.06563  0.04936  0.04239
## Cumulative Proportion  0.362  0.5541  0.6653  0.73599  0.80162  0.85098  0.89337
##                               PC8      PC9      PC10     PC11     PC12     PC13
## Standard deviation     0.59034 0.53748 0.5009  0.47517  0.41082  0.32152
## Proportion of Variance 0.02681 0.02222 0.0193  0.01737  0.01298  0.00795
## Cumulative Proportion  0.92018 0.94240 0.9617  0.97907  0.99205  1.00000

wn_pca$center

##          Alcohol      Malic_Acid           Ash       Alcalinity
## 13.0006180    2.3363483   2.3665169  19.4949438
##          Magnesium   Total_Phenols   Flavanoids Nonflavanoid
## 99.7415730    2.2951124   2.0292697  0.3618539
## Proanthocyanins Color_Intensity        Hue OD280/OD315
## 1.5908989    5.0580899   0.9574494  2.6116854
##          Proline
## 746.8932584


wn_pca$scale

##          Alcohol      Malic_Acid           Ash       Alcalinity
## 0.8118265    1.1171461   0.2743440  3.3395638
##          Magnesium   Total_Phenols   Flavanoids Nonflavanoid
## 14.2824835    0.6258510   0.9988587  0.1244533
## Proanthocyanins Color_Intensity        Hue OD280/OD315
## 0.5723589    2.3182859   0.2285716  0.7099904
##          Proline
## 314.9074743


```

```
plot(wn_pca)
```

wn_pca




```
biplot(wn_pca, scale = 0)
```



As from the above biplot, we can see that the feature opposite to Hue is Malic Acid which implies that they are inversely proportional which means if the value of one increases then the other decreases. This is further supported with the following correlation plot.

```
corrplot(cor(wn_data[, -1]), method="number")
```




```
cat("Correlation between Malic_Acid and Hue: ", cor(wn_data$Malic_Acid, wn_data$Hue))
```

```
## Correlation between Malic_Acid and Hue: -0.5612957
```

```
std_dev = wn_pca$sdev  
pr_var = std_dev^2  
prop_varex = pr_var/sum(pr_var)
```

```
plot(prop_varex, xlab = "Principal Component", ylab = "Proportion of Variance Explained", type = "b")
```



```
cat("Percentage of total variance explained by PC1 and PC2 = ",  
(prop_varex[1]+prop_varex[2])*100,"%")
```

```
## Percentage of total variance explained by PC1 and PC2 = 55.40634 %
```

2.3 Problem 3

```
library("factoextra")  
  
## Loading required package: ggplot2  
  
## Welcome! Related Books: `Practical Guide To Cluster Analysis in R` at  
https://goo.gl/13EFCZ  
  
data = data.frame(USArrests)  
states = row.names(data)  
states  
  
## [1] "Alabama"      "Alaska"       "Arizona"       "Arkansas"  
## [5] "California"   "Colorado"     "Connecticut"   "Delaware"  
## [9] "Florida"       "Georgia"      "Hawaii"        "Idaho"  
## [13] "Illinois"      "Indiana"      "Iowa"          "Kansas"  
## [17] "Kentucky"      "Louisiana"    "Maine"         "Maryland"  
## [21] "Massachusetts" "Michigan"     "Minnesota"    "Mississippi"  
## [25] "Missouri"      "Montana"      "Nebraska"     "Nevada"  
## [29] "New Hampshire" "New Jersey"   "New Mexico"   "New York"  
## [33] "North Carolina" "North Dakota" "Ohio"          "Oklahoma"  
## [37] "Oregon"        "Pennsylvania" "Rhode Island" "South Carolina"  
## [41] "South Dakota"   "Tennessee"   "Texas"        "Utah"
```

```

## [45] "Vermont"          "Virginia"          "Washington"        "West Virginia"
## [49] "Wisconsin"         "Wyoming"

names(data)

## [1] "Murder"   "Assault"   "UrbanPop"  "Rape"

# Before Scaling
head(data)

##           Murder Assault UrbanPop Rape
## Alabama     13.2    236      58 21.2
## Alaska      10.0    263      48 44.5
## Arizona      8.1    294      80 31.0
## Arkansas     8.8    190      50 19.5
## California    9.0    276      91 40.6
## Colorado      7.9    204      78 38.7

summary(data)

##           Murder       Assault      UrbanPop      Rape
## Min.   : 0.800   Min.   :45.0   Min.   :32.00   Min.   : 7.30
## 1st Qu.: 4.075   1st Qu.:109.0  1st Qu.:54.50  1st Qu.:15.07
## Median  : 7.250   Median :159.0   Median :66.00  Median :20.10
## Mean    : 7.788   Mean   :170.8   Mean   :65.54  Mean   :21.23
## 3rd Qu.:11.250   3rd Qu.:249.0  3rd Qu.:77.75  3rd Qu.:26.18
## Max.    :17.400   Max.   :337.0   Max.   :91.00  Max.   :46.00

```

As the means of the attributes differ, we need to apply Scaling.

```

# Scaling
new_data = scale(data)

# After Scaling
head(new_data)

##           Murder       Assault      UrbanPop      Rape
## Alabama     1.24256408  0.7828393 -0.5209066 -0.003416473
## Alaska      0.50786248  1.1068225 -1.2117642  2.484202941
## Arizona     0.07163341  1.4788032  0.9989801  1.042878388
## Arkansas    0.23234938  0.2308680 -1.0735927 -0.184916602
## California  0.27826823  1.2628144  1.7589234  2.067820292
## Colorado     0.02571456  0.3988593  0.8608085  1.864967207

summary(new_data)

##           Murder       Assault      UrbanPop      Rape
## Min.   :-1.6044   Min.   :-1.5090   Min.   :-2.31714  Min.   :-1.4874
## 1st Qu.:-0.8525   1st Qu.:-0.7411   1st Qu.:-0.76271  1st Qu.:-0.6574
## Median :-0.1235   Median :-0.1411   Median : 0.03178  Median :-0.1209
## Mean   : 0.0000   Mean   : 0.0000   Mean   : 0.00000  Mean   : 0.0000
## 3rd Qu.: 0.7949   3rd Qu.: 0.9388   3rd Qu.: 0.84354  3rd Qu.: 0.5277
## Max.   : 2.2069   Max.   : 1.9948   Max.   : 1.75892  Max.   : 2.6444

withinss = c()
i = 2
for (k in 2:10){
  set.seed(123)

```


```

model = kmeans(new_data, k, nstart = 25)
withinss[i] = model$tot.withinss
cat(" k = ",k,"Withinss = ",withinss[i],"\n")
i = i + 1
}

##  k = 2 Withinss = 102.8624
##  k = 3 Withinss = 78.32327
##  k = 4 Withinss = 56.40317
##  k = 5 Withinss = 48.9442
##  k = 6 Withinss = 42.83303
##  k = 7 Withinss = 38.25764
##  k = 8 Withinss = 34.44327
##  k = 9 Withinss = 30.17403
##  k = 10 Withinss = 26.18348

plot(withinss, xlab = "Values of K", main = "Values of within-cluster sum of squares for each value of k ", type = "b", col = "dark red")

```




As we can see there is a sharp decrease at $K = 4$, and from there onwards the withinss keeps on decreasing, thus we select $K = 4$ as optimal value.

```

model = kmeans(new_data, 4, nstart = 25)
fviz_cluster(model, data = new_data, ggtheme = theme_minimal(), main = " Optimal Clustering Plot for K = 4")

```



2.4 Problem 4

```

library(readr)
library(data.table)
library(corrplot)

## corrplot 0.84 loaded

library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:data.table':
## 
##     between, first, last

## The following objects are masked from 'package:stats':
## 
##     filter, lag

## The following objects are masked from 'package:base':
## 
##     intersect, setdiff, setequal, union

wq_URL = "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv"
wq_raw_data = fread(wq_URL, header = TRUE)
wq_data = wq_raw_data[,-12]

head(wq_data)

##   fixed acidity volatile acidity citric acid residual sugar chlorides
## 1:        7.0            0.27       0.36        20.7      0.045
## 2:        6.3            0.30       0.34          1.6      0.049
## 3:        8.1            0.28       0.40        6.9       0.050
## 4:        7.2            0.23       0.32        8.5       0.058

```

```

## 5:          7.2          0.23          0.32          8.5      0.058
## 6:          8.1          0.28          0.40          6.9      0.050
##   free sulfur dioxide total sulfur dioxide density    pH sulphates alcohol
## 1:             45            170  1.0010 3.00      0.45      8.8
## 2:             14            132  0.9940 3.30      0.49      9.5
## 3:             30             97  0.9951 3.26      0.44     10.1
## 4:             47            186  0.9956 3.19      0.40      9.9
## 5:             47            186  0.9956 3.19      0.40      9.9
## 6:             30             97  0.9951 3.26      0.44     10.1

summary(wq_data)

##   fixed acidity  volatile acidity  citric acid  residual sugar
## Min. : 3.800  Min. :0.0800  Min. :0.0000  Min. : 0.600
## 1st Qu.: 6.300 1st Qu.:0.2100 1st Qu.:0.2700 1st Qu.: 1.700
## Median : 6.800 Median :0.2600 Median :0.3200 Median : 5.200
## Mean   : 6.855 Mean   :0.2782 Mean   :0.3342 Mean   : 6.391
## 3rd Qu.: 7.300 3rd Qu.:0.3200 3rd Qu.:0.3900 3rd Qu.: 9.900
## Max.   :14.200 Max.   :1.1000 Max.   :1.6600 Max.   :65.800
##   chlorides    free sulfur dioxide total sulfur dioxide
## Min. :0.00900  Min. : 2.00  Min. : 9.0
## 1st Qu.:0.03600 1st Qu.: 23.00 1st Qu.:108.0
## Median :0.04300 Median : 34.00 Median :134.0
## Mean   :0.04577 Mean   : 35.31 Mean   :138.4
## 3rd Qu.:0.05000 3rd Qu.: 46.00 3rd Qu.:167.0
## Max.   :0.34600 Max.   :289.00 Max.   :440.0
##   density       pH        sulphates      alcohol
## Min. :0.9871  Min. :2.720  Min. :0.2200  Min. : 8.00
## 1st Qu.:0.9917 1st Qu.:3.090 1st Qu.:0.4100 1st Qu.: 9.50
## Median :0.9937 Median :3.180  Median :0.4700 Median :10.40
## Mean   :0.9940 Mean   :3.188  Mean   :0.4898 Mean   :10.51
## 3rd Qu.:0.9961 3rd Qu.:3.280 3rd Qu.:0.5500 3rd Qu.:11.40
## Max.   :1.0390 Max.   :3.820  Max.   :1.0800 Max.   :14.20

```

As the means of the attributes differ, we need to apply Scaling.

```

wq_data_scale = scale(wq_data)

head(wq_data_scale)

##   fixed acidity volatile acidity citric acid residual sugar   chlorides
## [1,] 0.1720794 -0.08176155 0.21325843 2.8210611 -0.03535139
## [2,] -0.6574340  0.21587359 0.04799622 -0.9446688 0.14773200
## [3,] 1.4756004  0.01745016 0.54378284 0.1002720 0.19350284
## [4,] 0.4090832 -0.47860841 -0.11726599 0.4157258 0.55966962
## [5,] 0.4090832 -0.47860841 -0.11726599 0.4157258 0.55966962
## [6,] 1.4756004  0.01745016 0.54378284 0.1002720 0.19350284
##   free sulfur dioxide total sulfur dioxide density      pH
## [1,] 0.5698734  0.7444890 2.331273996 -1.24679399
## [2,] -1.2528907 -0.1496693 -0.009153237 0.73995309
## [3,] -0.3121093 -0.9732363 0.358628185 0.47505348
## [4,] 0.6874711  1.1209768 0.525801559 0.01147916
## [5,] 0.6874711  1.1209768 0.525801559 0.01147916
## [6,] -0.3121093 -0.9732363 0.358628185 0.47505348
##   sulphates      alcohol
## [1,] -0.34914861 -1.3930102

```

```


## [2,]  0.00134171 -0.8241915
## [3,] -0.43677119 -0.3366326
## [4,] -0.78726151 -0.4991523
## [5,] -0.78726151 -0.4991523
## [6,] -0.43677119 -0.3366326

summary(wq_data_scale)

##   fixed acidity      volatile acidity      citric acid      residual sugar
## Min.   :-3.61998   Min.   :-1.9668   Min.   :-2.7615   Min.   :-1.1418
## 1st Qu.:-0.65743   1st Qu.:-0.6770   1st Qu.:-0.5304   1st Qu.:-0.9250
## Median :-0.06492   Median :-0.1810   Median :-0.1173   Median :-0.2349
## Mean    : 0.00000   Mean    : 0.0000   Mean    : 0.0000   Mean    : 0.0000
## 3rd Qu.: 0.52758   3rd Qu.: 0.4143   3rd Qu.: 0.4612   3rd Qu.: 0.6917
## Max.   : 8.70422   Max.   : 8.1528   Max.   :10.9553   Max.   :11.7129
##   chlorides      free sulfur dioxide      total sulfur dioxide
## Min.   :-1.6831   Min.   :-1.95848  Min.   :-3.0439
## 1st Qu.:-0.4473   1st Qu.:-0.72370  1st Qu.:-0.7144
## Median :-0.1269   Median :-0.07691  Median :-0.1026
## Mean    : 0.00000   Mean    : 0.00000  Mean    : 0.0000
## 3rd Qu.: 0.1935   3rd Qu.: 0.62867  3rd Qu.: 0.6739
## Max.   :13.7417   Max.   :14.91679  Max.   : 7.0977
##   density          pH          sulphates
## Min.   :-2.31280  Min.   :-3.10109  Min.   :-2.3645
## 1st Qu.:-0.77063  1st Qu.:-0.65077  1st Qu.:-0.6996
## Median :-0.09608  Median :-0.05475  Median :-0.1739
## Mean    : 0.00000  Mean    : 0.00000  Mean    : 0.0000
## 3rd Qu.: 0.69298  3rd Qu.: 0.60750  3rd Qu.: 0.5271
## Max.   :15.02976  Max.   : 4.18365  Max.   : 5.1711
##   alcohol
## Min.   :-2.04309
## 1st Qu.:-0.82419
## Median :-0.09285
## Mean    : 0.00000
## 3rd Qu.: 0.71974
## Max.   : 2.99502

hc.complete = hclust (dist(wq_data_scale), method = "complete")
plot(hc.complete, main =" Complete Linkage ", xlab = "")
abline(h =tail(hc.complete$height,1), col="red")

```



```

cat("Distance value where the two penultimate clusters merge =",  

tail(hc.complete$height,1),"\n")  

## Distance value where the two penultimate clusters merge = 27.73476  

hc.completeCut = cutree(hc.complete, 2)  

table(hc.completeCut)  

## hc.completeCut  

##   1   2  

## 4897   1  

wq_data2 = wq_data  

wq_data2$Clusters = hc.completeCut  

wq_data2 = dplyr::group_by(wq_data2,Clusters)  

a = dplyr::summarise_each(wq_data2,fun(mean))  

## Warning: funs() is soft deprecated as of dplyr 0.8.0  

## Please use a list of either functions or lambdas:  

##  

##  # Simple named list:  

##  list(mean = mean, median = median)  

##  

##  # Auto named with `tibble::lst()`:  

##  tibble::lst(mean, median)  

##  

##  # Using lambdas  

##  list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))  

## This warning is displayed once per session.  

print.data.frame(a)  

##   Clusters fixed acidity volatile acidity citric acid residual sugar  

## 1         1       6.854595      0.2781009    0.3341372     6.379283  

## 2         2       7.800000      0.9650000    0.6000000    65.800000  

##   chlorides free sulfur dioxide total sulfur dioxide density pH  

## 1 0.04576659           35.31366        138.3562 0.9940182 3.188225  

## 2 0.07400000            8.00000        160.0000 1.0389800 3.390000  

##   sulphates alcohol  

## 1  0.489806 10.51402  

## 2  0.690000 11.70000  

# Feature Means Difference  

abs(a[2,-1]-a[1,-1])  

##   fixed acidity volatile acidity citric acid residual sugar chlorides  

## 1  0.9454054     0.6868991  0.2658628  59.42072 0.02823341  

##   free sulfur dioxide total sulfur dioxide density pH sulphates  

## 1          27.31366        21.64376 0.0449618 0.2017746 0.200194  

##   alcohol  


## 1 1.185975  

hc.single = hclust (dist(wq_data_scale), method = "single")
plot(hc.single, main = " Single Linkage ", xlab = "")  

abline(h =tail(hc.single$height,1), col="red")

```

Single Linkage



```
hcclust (*, "single")
```

```
cat("Distance value where the two penultimate clusters merge =",  
tail(hc.single$height,1),"\n")  
  
## Distance value where the two penultimate clusters merge = 14.25323  
  
hc.singleCut = cutree(hc.single, 2)  
table(hc.singleCut)  
  
## hc.singleCut  
## 1 2  
## 4897 1  
  
wq_data3 = wq_data  
wq_data3$Clusters = hc.singleCut  
wq_data3 = dplyr::group_by(wq_data3,Clusters)  
b = dplyr::summarise_each(wq_data3,funs(mean))  
print.data.frame(b)  
  
## Clusters fixed acidity volatile acidity citric acid residual sugar  
## 1 1 6.854595 0.2781009 0.3341372 6.379283  
## 2 2 7.800000 0.9650000 0.6000000 65.800000  
## chlorides free sulfur dioxide total sulfur dioxide density pH  
## 1 0.04576659 35.31366 138.3562 0.9940182 3.188225  
## 2 0.07400000 8.00000 160.0000 1.0389800 3.390000  
## sulphates alcohol  
## 1 0.489806 10.51402  
## 2 0.690000 11.70000  
  
# Feature Means Difference  
abs(b[2,-1]-b[1,-1])  
  
## fixed acidity volatile acidity citric acid residual sugar chlorides  
## 1 0.9454054 0.6868991 0.2658628 59.42072 0.02823341  
## free sulfur dioxide total sulfur dioxide density pH sulphates  
## 1 27.31366 21.64376 0.0449618 0.2017746 0.200194
```

```
##    alcohol  
## 1 1.185975
```

As we can see, the Distance value where the two penultimate clusters merge for Complete Linkage is 27.73476 and the Distance value where the two penultimate clusters merge for Single Linkage is 14.25323. However, we are getting same number of observations in both cases, i.e. 1st Cluster contains 4897 observations and the 2nd Clusters contains 1 observation. Also, from Summary statistics, in both cases the feature **residual sugar** has the largest possible difference. Thus, complete linkage produces more balanced tree than single linkage.